

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import LogisticRegression
from sklearn.pipeline import Pipeline
from sklearn.metrics import accuracy_score, f1_score, classification_report
import joblib
```

```
!pip install --upgrade --force-reinstall transformers==4.44.2
!pip install --upgrade --force-reinstall accelerate datasets
```




```
import pandas as pd
from google.colab import files

# This opens a file picker in Colab → select your local CSV
uploaded = files.upload()

# Get the first uploaded filename
filename = list(uploaded.keys())[0]

# Read into pandas
df = pd.read_csv(filename)

# Quick sanity check
print("Rows:", len(df))
print("Columns:", df.columns.tolist())
print(df.head())
```

Choose files data.csv -none-any.whl.metadata (57 kB)
 data.csv (text/csv) 79825 bytes last modified: 23/09/2025 - 100% done
 Using data.csv for data.csv.9.0-py3-none-any.whl.metadata (10 kB)
 Collecting typing-extensions>=3.7.4.3 (from huggingface-hub<1.0,>=0.23.2->transformers==4.44.2)
 Using cached typing_extensions-4.15.0-py3-none-any.whl.metadata (3.3 kB)
 Traceback (most recent call last):
 0 File "/usr/local/lib/python3.12/dist-packages/pip/_internal/cli/base_command.py", line 179, in exc_logging_wrapp
 1 status = run_func(*args) Can we discuss pricing?? NEUTRAL
 2 Im excited to explore this further, plz send c... POSITIVE
 3 File "/usr/local/lib/python3.12/dist-packages/pip/_internal/cli/req_command.py", line 67, in wrapper
 4 return func(self, options, args) We not looking for new solutions. negative
 File "/usr/local/lib/python3.12/dist-packages/pip/_internal/commands/install.py", line 377, in run
 File "/usr/local/lib/python3.12/dist-packages/pip/_internal/commands/install.py", line 377, in run

```
df = df.rename(columns={"reply": "text", "label": "label"})
df["text"] = df["text"].astype(str).str.strip().str.replace(r"\s+", " ", regex=True)
df["label"] = df["label"].str.lower().str.strip()
```

```
# Quick check
print("Rows:", len(df))
print("Columns:", df.columns.tolist())
print("\nLabel distribution:\n", df["label"].value_counts())
df.head()
```

File "/usr/local/lib/python3.12/dist-packages/pip/_vendor/resolvelib/resolvers.py", line 239, in _attempt_to_pin
 ColumnNotFoundError: self.label updated criteria(candidate)
 File "/usr/local/lib/python3.12/dist-packages/pip/_vendor/resolvelib/resolvers.py", line 230, in _get_updated_cr
 label.add_to_criteria(criteria, requirement, parent=candidate)
 File "/usr/local/lib/python3.12/dist-packages/pip/_vendor/resolvelib/resolvers.py", line 173, in _add_to_criteri
 positive criterion.candidates:
 negative
 neutral 709
 NameError: local variable 'text' referenced before assignment
 File "/usr/local/lib/python3.12/dist-packages/pip/_vendor/resolvelib/structs.py", line 156, in __bool__
 return bool(self._sequence)
 text label
 0 Can we discuss pricing?? neutral
 1 Im excited to explore this further, plz send c... positive
 File "/usr/local/lib/python3.12/dist-packages/pip/_internal/resolution/resolvelib/found_candidates.py", line 174
 2 We not looking for new solutions. negative in self._incompatible_ids)
 File "/usr/local/lib/python3.12/dist-packages/pip/_internal/resolution/resolvelib/found_candidates.py", line 49,
 3 Lets, schedule a meeting to dive deeper positive
 File "/usr/local/lib/python3.12/dist-packages/pip/_internal/index/package_finder.py", line 884, in find_best_cand

Next steps: result = self.finder.find_best_candidate(
 Generate code with AI New interactive sheet

```
train_df, test_df = train_test_split(
    df, test_size=0.2, stratify=df['label'], random_state=42
)

print("Train size:", len(train_df), "Test size:", len(test_df))
```

File "/usr/local/lib/python3.12/dist-packages/pip/_internal/index/package_finder.py", line 792, in process_proje
 package_links = self.evaluate_links(
 File "/usr/local/lib/python3.12/dist-packages/pip/_internal/index/package_finder.py", line 792, in process_proje

```
pipeline = Pipeline([
    ("tfidf", TfidfVectorizer(
        lowercase=True,
        strip_accents="unicode",
        stop_words="english",
        ngram_range=(1, 2),
        max_features=20000
    )),
    ("clf", LogisticRegression(
```

```

        max_iter=200,
        solver="liblinear",
        random_state=42
    ))
])

param_grid = {"clf__C": [0.25, 0.5, 1.0, 2.0, 4.0]}
grid = GridSearchCV(
    pipeline, param_grid=param_grid,
    cv=5, n_jobs=-1, scoring="f1_macro", verbose=1
)

grid.fit(train_df['text'], train_df['label'])

```

Fitting 5 folds for each of 5 candidates, totalling 25 fits

```

GridSearchCV
└─ best_estimator_: Pipeline
    └─ TfidfVectorizer
    └─ LogisticRegression

```

File "/usr/lib/python3.12/dist-packages/pip/_internal/cli/base_command.py", line 100, in main
 /dist-packages/pip/_internal/cli/base_command.py", line 232, in _main
 /dist-packages/pip/_internal/cli/base_command.py", line 215, in exc_logging_wrapp
 cancelled by user")
 ng/_init_.py", line 1586, in critical
 , **kwargs)
 File "/usr/lib/python3.12/dist-packages/pip/_internal/cli/base_command.py", line 1684, in _log
 self.handle(record)

```

y_pred = grid.predict(test_df['text'])
print("Accuracy:", accuracy_score(test_df['label'], y_pred))
print("Macro F1:", f1_score(test_df['label'], y_pred, average="macro"))
print("\nClassification Report:\n", classification_report(test_df['label'], y_pred))

```

```

File "/usr/local/lib/python3.12/dist-packages/pip/_internal/commands/report.py", line 177, in emit
Accuracy: 0.98 0.98 0.98 0.98 0.98
Macro F1: 0.98 0.98 0.98 0.98 0.98
File "/usr/local/lib/python3.12/dist-packages/pip/_vendor/rich/console.py", line 1674, in print
renderables = self._collect_renderables(
Classification Report:
File "/usr/local/lib/python3.12/dist-packages/pip/_vendor/rich/console.py", line 1553, in _collect_renderables
check_text()
File "/usr/local/lib/python3.12/dist-packages/pip/_vendor/rich/console.py", line 1531, in check_text
a = text.join(text)
positive
File "/usr/local/lib/python3.12/dist-packages/pip/_vendor/rich/text.py", line 803, in join
766 text in iter_text():
macro avg 0.98 0.98 0.98 0.98 0.98
weighted avg 0.98 0.98 0.98 0.98 0.98
File "/usr/local/lib/python3.12/dist-packages/pip/_vendor/rich/text.py", line 790, in iter_text
for last, line in loop_last(lines):

```

```

import joblib
joblib.dump(grid.best_estimator_, "baseline_model.joblib")

```

```

['baseline_model.joblib']-----
KeyboardInterrupt Traceback (most recent call last)

```

```

from google.colab import files
files.download("baseline_model.joblib")

```

```

/usr/lib/python3.12/pathlib.py in stat(self, follow_symlinks)
838         os.stat() does.

```

```
!pip install -q transformers datasets accelerate evaluate
```

```
842 def lstat(self):
```

```
!pip install -q --upgrade transformers
```

```

import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split

import torch
from datasets import Dataset
from transformers import (
    AutoTokenizer,
    AutoModelForSequenceClassification,
    TrainingArguments,
    Trainer,
    DataCollatorWithPadding,
)

```

```
# Ensure columns are "text" and "label" with lowercase labels
df["label"] = df["label"].str.lower().str.strip()

# Train / validation split
train_df, val_df = train_test_split(
    df, test_size=0.2, stratify=df["label"], random_state=42
)

# Hugging Face dataset objects
train_ds = Dataset.from_pandas(train_df)
val_ds = Dataset.from_pandas(val_df)

# Label mapping
labels = ["negative", "neutral", "positive"]
label2id = {l: i for i, l in enumerate(labels)}
id2label = {i: l for l, i in label2id.items()}

def encode_labels(example):
    example["labels"] = label2id[example["label"]]
    return example

train_ds = train_ds.map(encode_labels)
val_ds = val_ds.map(encode_labels)
```

Map: 100% 1703/1703 [00:00<00:00, 4436.51 examples/s]

Map: 100% 426/426 [00:00<00:00, 4525.41 examples/s]

```
model_name = "distilbert-base-uncased"
tokenizer = AutoTokenizer.from_pretrained(model_name)

def tokenize(examples):
    return tokenizer(examples["text"], truncation=True, padding=False, max_length=256)

train_ds = train_ds.map(tokenize, batched=True)
val_ds = val_ds.map(tokenize, batched=True)

data_collator = DataCollatorWithPadding(tokenizer=tokenizer)
```

/usr/local/lib/python3.12/dist-packages/transformers/tokenization_utils_base.py:1601: FutureWarning: `clean_up_token` warnings.warn(

Map: 100% 1703/1703 [00:00<00:00, 7200.24 examples/s]

Map: 100% 426/426 [00:00<00:00, 4602.28 examples/s]

```
model = AutoModelForSequenceClassification.from_pretrained(
    model_name,
    num_labels=len(labels),
    id2label=id2label,
    label2id=label2id,
)
```

Some weights of DistilBertForSequenceClassification were not initialized from the model checkpoint at distilbert-base-uncased. You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

```
from sklearn.metrics import accuracy_score, f1_score
```

```
def compute_metrics(eval_pred):
    logits, labels = eval_pred
    preds = np.argmax(logits, axis=-1)
    acc = accuracy_score(labels, preds)
    f1 = f1_score(labels, preds, average="macro")
    return {"accuracy": acc, "f1": f1}
```

```
from transformers import TrainingArguments
training_args = TrainingArguments(
    output_dir="./distilbert-reply-clf",
    evaluation_strategy="epoch", # or "steps"
    save_strategy="epoch",
    learning_rate=2e-5,
    per_device_train_batch_size=16,
    per_device_eval_batch_size=16,
```

```

num_train_epochs=4,
weight_decay=0.01,
load_best_model_at_end=True,
metric_for_best_model="f1",
greater_is_better=True,
logging_dir="./logs",
logging_steps=20,
)

```

```

/usr/local/lib/python3.12/dist-packages/transformers/training_args.py:1525: FutureWarning: `evaluation_strategy` is
warnings.warn(

```

```
!pip install --upgrade transformers
```

```

import transformers
print(transformers.__version__)

```

```

from datasets import Dataset
from sklearn.model_selection import train_test_split

# 1. Normalize labels
df["label"] = df["label"].str.lower().str.strip()

# 2. Train/validation split
train_df, val_df = train_test_split(
    df, test_size=0.2, stratify=df["label"], random_state=42
)

# 3. Define mappings
labels = ["negative", "neutral", "positive"]
label2id = {l: i for i, l in enumerate(labels)}
id2label = {i: l for l, i in label2id.items()}

# 4. Convert pandas → Dataset, keeping only needed columns
train_ds = Dataset.from_pandas(train_df[["text", "label"]].reset_index(drop=True))
val_ds = Dataset.from_pandas(val_df[["text", "label"]].reset_index(drop=True))

# 5. Encode labels → integers
def encode_labels(example):
    return {"labels": label2id[example["label"]]}

train_ds = train_ds.map(encode_labels)
val_ds = val_ds.map(encode_labels)

# 6. Tokenizer
from transformers import AutoTokenizer
model_name = "distilbert-base-uncased"
tokenizer = AutoTokenizer.from_pretrained(model_name)

def tokenize(batch):
    return tokenizer(batch["text"], truncation=True, padding=True, max_length=256)

train_ds = train_ds.map(tokenize, batched=True)
val_ds = val_ds.map(tokenize, batched=True)

# 7. Remove original string label column
train_ds = train_ds.remove_columns(["label"])
val_ds = val_ds.remove_columns(["label"])

# 8. Verify
print(train_ds[0])

```

```
Map: 100% 1703/1703 [00:00<00:00, 9920.73 examples/s]
```

```
Map: 100% 426/426 [00:00<00:00, 4857.16 examples/s]
```

```

/usr/local/lib/python3.12/dist-packages/transformers/tokenization_utils_base.py:1601: FutureWarning: `clean_up_token
warnings.warn(

```

```
Map: 100% 1703/1703 [00:00<00:00, 6723.83 examples/s]
```

```
Map: 100% 426/426 [00:00<00:00, 3659.75 examples/s]
```

```
{'text': 'Please share the details, I'm interested.', 'labels': 2, 'input_ids': [101, 3531, 3745, 1996, 4751, 1010,
```

```

from transformers import AutoModelForSequenceClassification

model = AutoModelForSequenceClassification.from_pretrained(
    "distilbert-base-uncased",

```

```

num_labels=len(labels),
id2label=id2label,
label2id=label2id,
)

```

Some weights of DistilBertForSequenceClassification were not initialized from the model checkpoint at distilbert-base-uncased. You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

```

from transformers import TrainingArguments

training_args = TrainingArguments(
    output_dir="./distilbert-reply-clf",
    eval_strategy="epoch",           # use "eval_strategy" (new name)
    save_strategy="epoch",
    learning_rate=2e-5,
    per_device_train_batch_size=16,
    per_device_eval_batch_size=16,
    num_train_epochs=4,
    weight_decay=0.01,
    load_best_model_at_end=True,
    metric_for_best_model="f1",
    greater_is_better=True,
    logging_dir="./logs",
    logging_steps=20,
    report_to="none",               # ✅ turn off wandb prompt
)

```

```

from sklearn.metrics import accuracy_score, f1_score

def compute_metrics(eval_pred):
    logits, labels = eval_pred
    preds = logits.argmax(axis=-1)
    acc = accuracy_score(labels, preds)
    f1 = f1_score(labels, preds, average="macro")
    return {"accuracy": acc, "f1": f1}

```

```

from transformers import Trainer, DataCollatorWithPadding

data_collator = DataCollatorWithPadding(tokenizer=tokenizer)

trainer = Trainer(
    model=model,
    args=training_args,
    train_dataset=train_ds,
    eval_dataset=val_ds,
    tokenizer=tokenizer,
    data_collator=data_collator,
    compute_metrics=compute_metrics,
)

```

```
trainer.train()
```

/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py:666: UserWarning: 'pin_memory' argument is set to False but pin_memory_device is specified. This could lead to unexpected behavior. Please set pin_memory=True to silence this warning.
warnings.warn(warn_msg)

[428/428 18:11, Epoch 4/4]

Epoch	Training Loss	Validation Loss	Accuracy	F1
1	0.024900	0.025807	0.995305	0.995305
2	0.006000	0.004189	1.000000	1.000000
3	0.003800	0.002724	1.000000	1.000000
4	0.003100	0.002383	1.000000	1.000000

/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py:666: UserWarning: 'pin_memory' argument is set to False but pin_memory_device is specified. This could lead to unexpected behavior. Please set pin_memory=True to silence this warning.
warnings.warn(warn_msg)

/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py:666: UserWarning: 'pin_memory' argument is set to False but pin_memory_device is specified. This could lead to unexpected behavior. Please set pin_memory=True to silence this warning.
warnings.warn(warn_msg)

/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py:666: UserWarning: 'pin_memory' argument is set to False but pin_memory_device is specified. This could lead to unexpected behavior. Please set pin_memory=True to silence this warning.
warnings.warn(warn_msg)

/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py:666: UserWarning: 'pin_memory' argument is set to False but pin_memory_device is specified. This could lead to unexpected behavior. Please set pin_memory=True to silence this warning.
warnings.warn(warn_msg)

TrainOutput(global_step=428, training_loss=0.07845190939467366, metrics={'train_runtime': 1104.7997, 'train_samples_per_second': 6.166, 'train_steps_per_second': 0.387, 'total_flos': 24674562830736.0, 'train_loss':